

Component Technology - Middleware

110011264

School of Computer Science
University of St Andrews

110002255

School of Computer Science
University of St Andrews

ABSTRACT

Abstract

REPORT

Planning

Initially, the team implemented a simple server. The server receives heartbeats, in the form of POST request, from the producers. Each heartbeat consisted of an unique identifier of the mobile phone device and its current location. The sender advertises itself as a producer to the server. When a consumer asks the server for sensor data from that particular producer, the server will send a request to the producer, and then pass on the results back to the producer. However, this approach requires the producers listening on some ports to handle data requests from the server. As a result, the consumer will spend significant time waiting for the server to retrieve data from the producer. The solution will create too much overhead on the server side, and it is unlikely that the server can deal with increasing consumer requests.

Later on, the team arrived to a more elegant solution. On each heartbeat, the producer sends sensor data to the server. The server caches the data, and flushes them to a persistent database on occasion. Now, when a consumer requests sensor data from a particular producer, the server can respond quickly, since the sensor data are cached. A more detailed description of the overall setup is explained in the next

Instructions

See "README.md" in the project directory for directions on how to run the servers, producers, and consumers.

Design

System components

The final system architecture at the time of our submission consists of:

- 1 Load balancer
- N Workers
- X Producers
- Y Consumers

A load balancer is the primary component of the system. It manages every other system component and balances the server load. A worker is a separate unit that is responsible for data retrieval and storage. A producer generates sensor data, and a consumer requests those data. An example of a producer would be a mobile phone, and an example of a consumer would be a web service that uses location data to recommend music to its users.

Load balancer and workers

The load balancer is the entry point of the entire system. It is initialised to be listening on a specified port. It supports a number of RESTful API calls that allows other entities to interact with its system components, such as a worker. The worker is initialised with the IP address and the port number of the load balancer, so that it can register itself as a worker for the load balancer via with the `/balancer/port` API call. The load balancer processes the registration request, and stores the worker's IP address in memory. The load balancer keeps track of every worker, and will distribute the server load equally, using the round robin method. Every worker must go through the above registration process in order to let itself known to the load balancer, whom will redirect producers and consumers to the worker. The worker also listens on a specified port, for handling heartbeats (described below) and sensor data requests.

Load balancer and producers

The producers need to go through the balancer in order to get a worker assigned to them. The balancer uses the list of online workers to assign a worker. A round robin approach is used to ensure that the load is distributed evenly to all workers and is almost the same at every given time for every worker. An API call `/register/{int:id}` is made by the producer to the balancer with its ID. The balancer keeps track of the relationship between producers and workers which is used when the consumer requests are handled. When a worker is found, the producer gets redirected to it and the communication between the producer and balancer is over.

A typical use case: A producer is started by providing a balancer address and producer ID. The finds a worker for the producer or returns a message that no workers are available. If a worker is available the corresponding address is returned. All errors regarding connection loss or interruption are handled.

Workers and producers: heartbeats

After a producer receives its assigned worker address it starts sending heartbeats (periodic messages). They contain the following information:

Code Fragment 1. Heartbeat format

A typical use case: A producer receives a worker address from the balancer and uses it to start making heartbeats to the worker, providing data and information. If the worker happens to go down, a retry mechanism triggers and the producers tries to re-register with the balancer. If workers are available the producer gets a new worker and continues heartbeats. All errors regarding connection loss or interruption are handled.

Load balancer and consumers: redirection

A typical use case: A consumer makes a GET request to the balancer. The balancer checks if the ID of the specified producer is being handled by some of the workers. If so the consumer get redirected to the appropriate worker which will

Workers and consumers

A typical use case: A consumer get redirected to a worker. The worker checks if live data for the ID is available in the cache. If not it makes a query to the database. The data is returned to the consumer.

Becuase

Because

Premature optimization is the root of all evil.

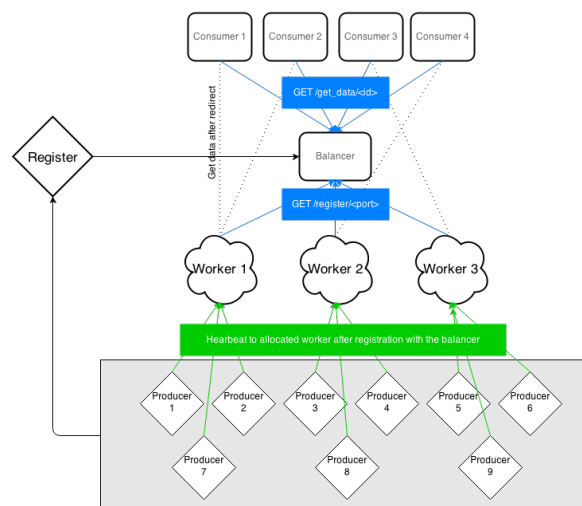


Figure 1. The final system architecture

Testing

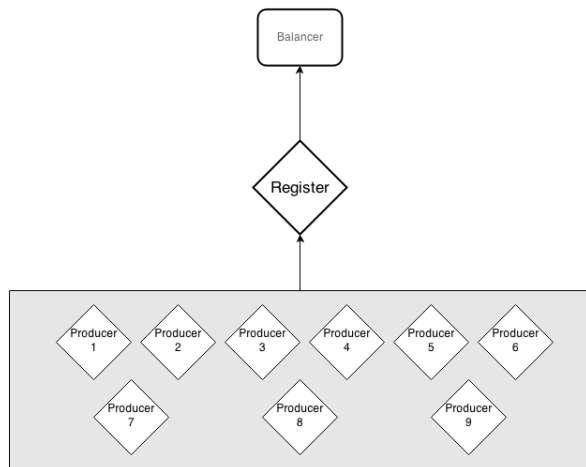


Figure 2. A producer first registers with the balancer and is assigned a worker

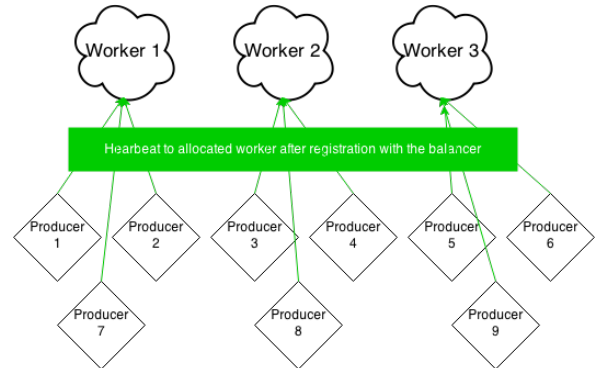


Figure 5. Producers send heartbeats to their assigned workers.

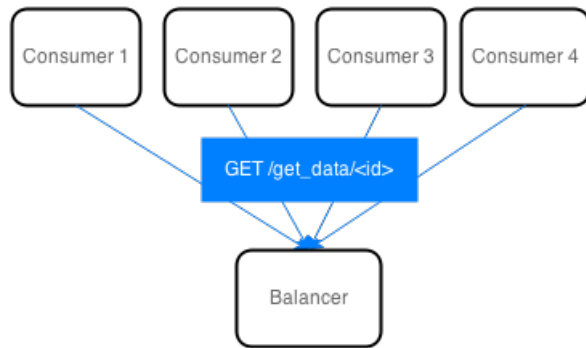


Figure 3. Consumers request data from the load balancer




Figure 4. Workers need to register with the balancer before they can start accepting work.




Figure 6. Cons to worker